

Anthropic’s J-Lens Is Not a Mind Reader

The Real, The Hype, and the Impact

Ben Mancke

July 14, 2026

TL;DR

Anthropic’s recent J-space work is scientifically interesting, but not for the reasons emphasized in much of the public discussion. Anthropic’s work presents a precomputed layer-wise averaged Jacobian: a first-order approximation of how perturbations to an intermediate residual-stream state affect downstream output-facing states. Composed with the model’s unembedding, this gives a readout of which vocabulary tokens an activation is, on average, disposed to make the model verbalize. The surprising result is that this static averaged tangent operator remains semantically coherent enough to expose and causally manipulate some verbalizable intermediate representations. That is useful. What it isn’t is evidence of consciousness, subjective experience, or an inner mind. And importantly, it doesn’t get us over the finish line of a robust, complete mechanistic theory of transformer dynamics.

Why I Care About This Result

The majority of my free time for over a year has been studying Jacobian and tangent-space effects inside transformer models: how perturbations to hidden states propagate through layers, how those effects differ by mechanism, and how local learned structure is preserved or distorted during generation. Naturally, Anthropic’s J-lens immediately caught my attention.

Applying an averaged downstream Jacobian to a hidden state and then unembedding isn’t wildly exotic. Anyone with some solid math foundations should be asking about Jacobians and Hessians once they get their AI feet wet. The interesting part is that an averaged downstream Jacobian appears to preserve coherent information at all. A transformer is nonlinear, context-dependent, attention-pattern-dependent, position-dependent, and mediated by many interacting components. I would have expected averaging tangent effects across contexts to wash out the dynamics we’re trying to understand. Anthropic’s result suggests that, at least in some cases, it does not. That’s a useful mathematical object that should be a significant character in the story of the Transformer Jacobian rabbit-hole.

The Real

Anthropic’s Jacobian

Let $h_{\ell,t}$ denote the residual-stream activation at layer ℓ and token position t . A perturbation to this activation affects later final-layer (denoted L) residual states $h_{L,t'}$, for positions t' . To first order, this effect is described by a Jacobian $\frac{\partial h_{L,t'}}{\partial h_{\ell,t}}$.

Anthropic’s J-lens averages this object over source positions, future positions, and a corpus of prompts, producing one static matrix per layer:

$$\mathcal{J}_\ell = \mathbb{E}_{t, t' \geq t, p} \left[\frac{\partial h_{L, t'}^{(p)}}{\partial h_{\ell, t}^{(p)}} \right].$$

Here p indexes prompts. The result is a single $d_{\text{model}} \times d_{\text{model}}$ matrix for each layer ℓ . Applying the lens to a hidden state is then approximately

$$\text{JLens}_\ell(h_\ell) = \text{softmax}(W_U \text{norm}(\mathcal{J}_\ell h_\ell)).$$

In plain language, the J-lens asks:

If a hidden state were propagated through an average linear approximation of the remaining transformer computation, which output words would it tend to make more likely?

That is a useful question. But it is also a much more concrete question than “what is the model thinking?”

What *Is* a Big Deal

Anthropic’s averaged downstream Jacobian $\mathbb{E} \left[\frac{\partial h_L}{\partial h_\ell} \right]$ appears to recover, in many but not all cases, a meaningful output-facing coordinate system. That is absolutely nontrivial. The exact downstream map from $h_{\ell, t}$ to later outputs depends on the current prompt, token position, attention patterns, layer normalization state, residual-stream scale, and local nonlinear activation regime. I would have expected much of that structure to wash out under averaging. But Anthropic finds that an averaged first-order map can expose verbalizable internal content and support causal interventions. In some examples, changing the J-lens coordinate corresponding to one latent concept changes the model’s downstream answer in the expected way. Well, ain’t that somethin’?

Here’s the punchline:

- The mathematical object is real: an averaged downstream Jacobian can expose a sparse, verbalizable, output-facing component of residual-stream activity. That’s a fair claim.
- Anthropic provides evidence that this component can be causally and functionally important in selected regimes. No pushback here.
- The title presents this structure as the model’s “global workspace,” and the press release runs with it. Sorry, but I’m not yet convinced. In Section 1.3 the authors make a narrower claim of *some* workspace properties. They decline to claim the model reproduces the full workspace architecture. I’ll take them at their word.

I’m genuinely surprised and impressed by what J-space achieves. That being said, a reusable, verbally accessible readout geometry with some workspace-like structure is evidence that something workspace-like exists in the model. But treating that as *the* sought-after global workspace itself is an architectural jump. And when the technical details acknowledge a narrower access-style claim while the title and press frame it as the full monty global-workspace, we get mixed messaging that feeds cascading misunderstandings in the telephone-game media landscape.

The Hype

Brass Tacks: It Is a Linearized Output Readout

The basic lens operation is not mysterious:

$$h_\ell \mapsto \mathcal{J}_\ell h_\ell \mapsto \text{softmax}(W_U \text{norm}(\mathcal{J}_\ell h_\ell)).$$

This is a linearized output-sensitivity readout. It is a principled refinement of the logit lens, where the downstream transformer computation is approximated by a learned or computed layer-wise map rather than ignored.

That does not make it unimportant. Simple objects can be powerful. But the simplicity matters when interpreting the result. The object mathematically licenses a linearized output-sensitivity readout: an averaged downstream tangent map composed with an unembedding. That is already interesting, but it is not the same thing as a full cognitive — or “cognoidal” — architecture; a complete account of local model dynamics; or evidence of internal “thinking” in any conventional or serious sense.

Those of us who slugged through a math undergrad, or learned this stuff in the wild, should remember what a Jacobian is: the coordinate representation of a first-order total derivative. It gives a local linear approximation to a differentiable map. The J-lens is built from a layer-indexed average of that object, and J-space is then built from the resulting token-indexed directions. What the result shows is that this averaged derivative can produce meaningful output-facing approximations of downstream transformer computation. Surviving averaging is the surprising part. The fact that a derivative locally approximates the map it differentiates is the part we should already expect. J-space should not be represented as the model’s subconscious mind or as “what the model is thinking about.” Claims like that are science fiction, no matter who is making them. Anyone who wants to define “thinking” broadly enough to make it land here is invited to apply that definition consistently elsewhere.

J-space is a token-indexed readout geometry over residual-stream activity under a genuinely clever linear approximation. Also, keep in mind that J-space is not a full linear subspace in the conventional sense. The J-lens vectors form an overcomplete vocabulary-indexed frame. In practice, J-space is the sparse nonnegative bundle of J-lens token directions that best approximates the relevant component of the hidden state. That is a genuinely neat construction: an overcomplete vocabulary-indexed frame, constrained into a sparse positive readout geometry. Well played, Anthropic. Touché.

I’d call it a constrained sparse readout geometry of the hidden state. That is already plenty interesting. Inflating it into “what the model is thinking about” is handwavy at best, and memetically toxic at worst. Good science requires that the object be considered in terms of theorems implied by its mathematical definition, as well as what is learned from experiments.

Transport-Like, But Not a Full Transport Theory

I think it’s fair to call Anthropic’s $\mathcal{J}_\ell$ “transportoidal”. It approximates how a perturbation at layer ℓ would propagate into final-layer, output-facing coordinates: layer-wise average of downstream tangent transport. But it should not be confused with a mathematically complete differential-transport mechanism. We still need more ingredients and prep work before we can cook that dish.

A full differential-transport account would need to specify local, context-dependent maps; how those maps compose along the actual transformer path; what frame or metric makes comparisons meaningful; and how layer depth and position-indexed attention, feature activation, residual,

normalization, and interaction terms separately contribute to the total effect. The J-lens collapses all of that into one static averaged matrix per layer. They show that a layer-indexed averaged tangent map is enough to recover real structure, but a mature theory needs a full set of mathematical abstractions that result in a complete differential accounting of the model.

To be fair, that’s a long road and there are a lot of left turns to avoid. Nevertheless, this is what we need to require if we are going to really get out of the black box.

In other words, Anthropic studies something like

$$\mathcal{J}_\ell^{\text{avg}} = \mathbb{E} [D\Phi_{\ell \rightarrow L}],$$

where $\Phi_{\ell \rightarrow L}$ is the downstream transformer map from layer ℓ to the final layer. But a local mechanistic account would ask about objects closer to a local prompt-specific tangent map $J_\ell^{(p,t)}$

$$J_\ell^{(p,t)} = D\Phi_{\ell \rightarrow L}^{(p,t)},$$

and then ask how this object separates by mechanism, position, head, layer, and interaction pathway.

A static averaged Jacobian may reveal reusable background geometry. But it doesn’t explain how a specific prompt’s computation unfolds through the transformer. It’s a very cool lens, but probably a lower-dimensional projection of the target phenomenology.

What the Public Framing Gets Wrong

The public discussion around J-space has cascaded within days toward consciousness speculation. Disappointing but not surprising. The public-facing language surrounding the work uses terms such as “thinking,” “mind,” and “what the model is thinking but not saying.” Often, this happens in spite of the authors’ own disclaimers.

The problem is that the disclaimer does not fully undo the framing. The minute J-lens is described as a window into a model’s thoughts, casual readers are invited (and I fear encouraged) to treat the result as evidence about inner experience rather than as a readout of residual-stream content.

The scientific result does not require that baggage. In fact it hinders the progress and clouds understanding. A hidden activation can be verbalizable, causally patchable, useful for reasoning, and alignment-relevant without being described (even implicitly) as phenomenological evidence of consciousness. You can just say “The J-lens reads hidden states through an averaged downstream linearization into output-token space,” and leave it at that. That’s plenty enough flex.

The Impact

Significant Fair Claims:

- Some transformer hidden states contain output-verbalizable intermediate content.
- A static averaged downstream Jacobian can recover part of that content.
- Some of the recovered directions are causally relevant: intervening on them can change downstream behavior.
- The output-verbalizable component appears to be selective rather than identical to the entire residual stream.

- The J-space may have workspace-like functional properties in selected layers and selected tasks.

What Goes Too Far:

- Calling this anything on the order of evidence of consciousness;
- Treating a sparse verbalizable readout geometry as if it were the model’s complete global workspace;
- Assuming important reasoning is necessarily readable in J-space;
- Treating J-space as a mechanistically separated accounting of transformer dynamics;
- Assuming layer-indexed static averaged Jacobians reliably account for important local dynamics.

I don’t mean to be pedantic. But *differentiating* between a useful interpretability method and exaggerations is something we should insist on. This stuff exists at the peak intersection of pop culture and high impact implementations. Due diligence requires we don’t let it ride uncorrected.

The Next Serious Questions

The most interesting follow-up is not whether one can find more evocative words in the J-lens. The interesting question is what survives if the averaged tangent object is decomposed.

For example:

$$\mathcal{J}_\ell^{\text{total}} = \mathbb{E}[D\Phi_{\ell \rightarrow L}]$$

could be compared against mechanism-conditioned objects:

$$\mathcal{J}_\ell^{\text{attn}}, \quad \mathcal{J}_\ell^{\text{mlp}}, \quad \mathcal{J}_\ell^{\text{interaction}},$$

or against position-conditioned objects:

$$\mathcal{J}_\ell^{t \rightarrow t'} = \mathbb{E}_p \left[\frac{\partial h_{L,t'}^{(p)}}{\partial h_{\ell,t}^{(p)}} \right].$$

One could also compare the static average map to the local prompt-specific tangent map:

$$R_\ell^{(p,t)} = J_\ell^{(p,t)} - \mathcal{J}_\ell^{\text{avg}}.$$

One of the first thoughts I had was to ask what we see when inspecting residual $R_\ell^{(p,t)}$. This may be where much of the local computation lives: task-specific routing, mechanism-specific misrepresentation, context-specific semantic transport, or failure-mode onset.

This is the direction that interests me most. You could say it rhymes with the last year of my “work outside of work”. To keep it brief and avoid giving away the sauce: mechanism-separated tangent effects on fields interpolated from the model’s learned structure; a mathematically grounded first-principles phenomenological approach. And along with it, promising results that include reproducible behavior prediction signals across models.

If a layer-wise averaged Jacobian already recovers coherent verbalizable structure, then richer tangent objects may expose not merely what an activation can be made to say, but how learned structure is impacted by the transformer.

Conclusion

The J-space work is absolutely valuable, but the value should be accurately understood. The big deal is not that they unembed a hidden state after applying a layer-averaged downstream Jacobian. The big deal is that the averaged downstream tangent map remains coherent enough to expose and manipulate some internal representations. That fact has legs.

What needs to stop is the public-facing “mind reading” narrative. The objective is trained token prediction; the mechanism that serves it is structured, nontrivial, and worth studying. But none of that structure is evidence of inner experience. And as clever mathematical tricks go, J-space is one of the coolest we have seen in a minute. I just don’t think it’s the only character in the story of Differential Transformer Phenomenology.

References

Anthropic / Transformer Circuits, “Verbalizable Representations Form a Global Workspace in Language Models,” 2026.

<https://transformer-circuits.pub/2026/workspace/index.html>

Anthropic, “A global workspace in language models,” 2026.

<https://www.anthropic.com/research/global-workspace>